

LAB: String Patterns, Sorting & Grouping

The practice problems for this Lab will provide hands on experience with string patterns, sorting result sets and grouping result sets. You will also learn how to run SQL scripts to create several tables at once, as well as how to load data into tables from .csv files.

HR Database

We will be working on a sample HR database for this Lab. This HR database schema consists of 5 tables called EMPLOYEES, JOB_HISTORY, JOBS, DEPARTMENTS and LOCATIONS. Each table has a few rows of sample data. The following diagram shows the tables for the HR database.

SAMPLE HR DATABASE TABLES

EMPLOYEES

EMP_ID	F_NAME	L_NAME	SSN	B_DATE	SEX	ADDRESS	JOB_ID	SALARY	MANAGER_ID	DEP_ID
E1001	John	Thomas	123456	1976-01-09	M	5631 Rice, OakPark,IL	100	100000	30001	2
E1002	Alice	James	123457	1972-07-31	F	980 Berry Ln, Elgin,IL	200	80000	30002	5
E1003	Steve	Wells	123458	1980-08-10	M	291 Springs, Gary,IL	300	50000	30002	5

JOB_HISTORY

EMPL_ID	START_DATE	JOBS_ID	DEPT_ID
E1001	2000-01-30	100	2
E1002	2010-08-16	200	5
E1003	2016-08-10	300	5

JOBS

JOB_IDENT	JOB_TITLE	MIN_SALARY	MAX_SALARY
100	Sr. Architect	60000	100000
200	Sr.SoftwareDeveloper	60000	80000
300	Jr.SoftwareDeveloper	40000	60000

DEPARTMENTS

DEPT_ID_DEP	DEP_NAME	MANAGER_ID	LOC_ID
2	Architect Group	30001	L0001
5	Software Development	30002	L0002
7	Design Team	30003	L0003
5	Software	30004	L0004

LOCATIONS

LOCT_ID	DEP_ID_LOC
L0001	2
L0002	5
L0003	7

To complete this lab you will utilize Db2 database service on IBM Cloud as you did for the previous lab. There are three parts to this lab:

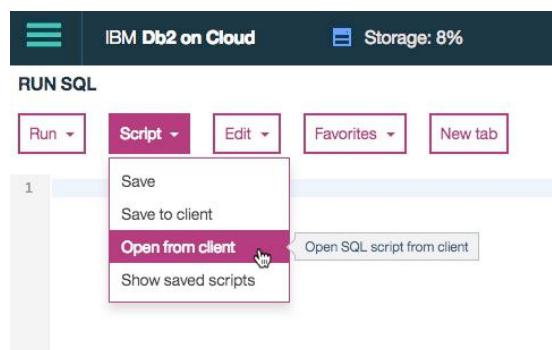
- I. Creating tables
- II. Loading data into tables
- III. Composing and running queries

Part I: CREATING TABLES

If you do not yet have access to Db2 on IBM Cloud, please refer to Lab Instructions in the Module/Week 1.

Rather than create each table manually by typing the DDL commands in the SQL editor, you will execute a script containing the create table commands for all the tables. Following step by step instructions are provided to perform this:

- 1) Download the script file “ Script_Create_Tables.sql” provided on the Lab page
- 2) Login to IBM Cloud and go to the Dashboard:
<https://console.bluemix.net/dashboard/apps> where you can find the Db2 service that you created in a previous Lab. Click on the service. Next, open the Db2 Console by clicking on "Open Console" button. Go to the Run SQL page. The Run SQL tool enables you to run DDL and SQL statements.
- 3) Click on the Folder icon to "Open script". Locate the file Script_Create_Tables.sql that you downloaded to your computer earlier and open it.



- 4) Once the statements are in the SQL Editor tool , you can run the queries against the database by selecting the option “Run” then "Run All".

A screenshot of the IBM Db2 on Cloud Run SQL interface. At the top, there are buttons for 'Run', 'Script', 'Edit', 'Favorites', and 'New tab'. The 'Run' button is highlighted with a red box and a mouse cursor. Below the buttons, there's a code editor area with a vertical line of numbers on the left (1, 2, 3, 4, 5, 6, 7, 8, 9). Lines 1 through 4 are dashed and contain the text '--DDL statement for table 'HR' database--'. Line 5 starts with 'CREATE TABLE EMPLOYEES (' and continues with column definitions: 'EMP_ID CHAR(9) NOT NULL', 'F_NAME VARCHAR(15) NOT NULL', 'L_NAME VARCHAR(15) NOT NULL', and 'SSN CHAR(9)', each on a new line. Lines 6 through 9 are dashed.

5) At the bottom of the screen you will see a Result section. Clicking on a query will show the execution details of the job - whether it ran successfully, or had any errors or warnings. Ensure your queries ran successfully and created all the tables.

The screenshot shows a 'Result' tab interface. On the left, a list of queries is shown, each with a green checkmark indicating success. The queries include creating tables for JOB_HISTORY, JOBS, DEPARTMENTS, and LOCATIONS. To the right, a 'Log' section displays the message 'The statement ran successfully.' Below the log, the 'SQL statement' is shown as 'CREATE TABLE LOCATIONS (...)'. Further down, the 'Execution log' section shows 'Run time: 0:103 s' and 'Status: SUCCEEDED'.

6) You can delete the Jobs history by clicking on the Delete icon as shown below.

The screenshot shows a 'Jobs' page with a 'Finished successfully' section. It lists five completed tasks under 'All (5)'. A red circle highlights the delete icon (a trash bin) located next to the 'Failed (0)' button.

7) Now you can look at the tables you created. Navigate to the three bar menu icon and select "Explore". Select the correct Schema to see the newly created tables.

The screenshot shows the 'Explore' interface of IBM Db2 on Cloud. The top navigation bar includes 'IBM Db2 on Cloud', 'Storage: 27%', 'Discover', 'Bell', and a user icon. The main area shows the 'EXPLORE' section. On the left, the 'Schema' list shows 'QCM54853' (circled in red), 'ERRORSCHEMA Sample', and 'ST_INFORMTN_SCHEMA Sample'. In the center, the 'Table' list shows 'DEPARTMENTS' (circled in red), 'EMPLOYEES', 'INSTRUCTOR', 'JOBS', 'JOB_HISTORY', and 'LOCATIONS'. On the right, the 'Table definition' panel for 'DEPARTMENTS' shows the following columns:

COLUMN	DATA TYPE	NULLABLE
NAME		
DEPT_ID_...	CHARA...	N
DEP_NAME	VARCHAR	Y
MANAGE...	CHARA...	Y
LOC_ID	CHARA...	Y

Part II: LOADING DATA

Now let us see how data can be loaded into Db2. We could manually insert each row into the table one by one but that would take a long time. Instead, Db2 (and almost every other database) allows you to Load data from csv files.

Please follow the steps below which explains the process of loading data into the tables we created earlier.

1) Download the 5 required data source files:

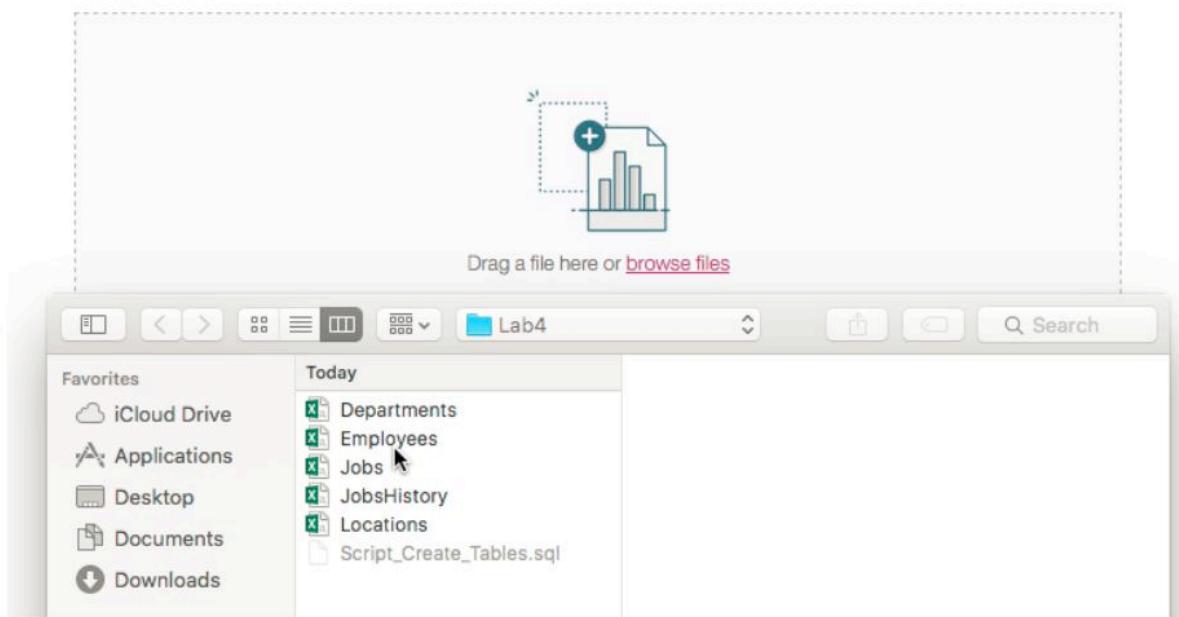
("Employees.csv", "Departments.csv", "Jobs.csv", "JobsHistory.csv", "Locations.csv") to your computer:

2) First let us learn how to load data into the Employees table that we created earlier. From the 3 bar menu icon, navigate to the "Load" page and ensure "My Computer" is selected as the source. Click on the "browse files" link.

The screenshot shows the IBM Db2 on Cloud interface with the 'LOAD' button highlighted. The 'Source' tab is active, showing a list of options: 'My Computer' (selected and circled in red), 'Amazon S3', and 'SoftLayer Swift'. Below this, there's a note about 'Netezza and large CSV file migrations' and a 'Lift' option. To the right, a 'File selection' panel has a dashed border. It contains a 'Drag a file here or browse files' area with a plus sign icon, also circled in red. At the bottom of the panel are 'High-speed loads powered by Aspera' and a 'Next' button.

3) Choose the file "Employees.csv" that you downloaded to your computer and click "Open".

File selection



4) Once the File is selected click "Next" in the bottom right corner.

LOAD

5) Select the schema for your Userid. It will show all the tables that have been created in this Schema previously, including the Employees table. Select the EMPLOYEES table.

You are loading the file **Employees.csv** into **QCM54853.EMPLOYEES**

Select a load target

Schema	Table	Table definition
Find a schema <input type="text"/>	Find a table in QCM54853 <input type="text"/>	EMPLOYEES Approximate 0 rows (0 KB) Updated on 5/2/2018 at 1:09:31 PM <input checked="" type="radio"/> Append new data <input type="radio"/> Overwrite table with new data
QCM54853 <input checked="" type="radio"/>	DEPARTMENTS	COLUMN DATA TYPE NULLABLE
ERRORSCHEMA <i>Sample</i>	EMPLOYEES <input checked="" type="radio"/>	NAME
ST_INFORMTN_SCHEMA <i>Sample</i>	INSTRUCTOR	EMP_ID CHARA...
	JOB_HISTORY	F_NAME VARCHAR

[Back](#) [Next](#)

- 6) Since our source data files do not contain any rows with column labels, turn off the setting for "Header in first row". Also, click on the down arrow next to "Date format" and choose "MM/DD/YYYY" since that is how the date is formatted in our source file.

You are loading the file **Employees.csv** into **QCM54853.EMPLOYEES**

Code page (character encoding): 1208 (UTF-8) <input type="button" value="▼"/>	Separator: , <input type="button" value="▼"/>	Header in first row: <input checked="" type="checkbox"/> <input type="button" value="Time & date format: ?"/>				
Date format: MM/DD/YYYY <input type="button" value="▼"/>	Time format: HH:MM:SS <input type="button" value="▼"/>	Timestamp format: YYYY-MM-DD HH:MM:SS <input type="button" value="▼"/>				
EMP_ID CHARACTER	F_NAME VARCHAR	L_NAME VARCHAR	SSN CHARACTER	B_DATE DATE	SEX CHARACTER	ADDRESS VARCHAR
1 E1001	John	Thomas	123456	01/09/1976	M	"5631 Rice
2 E1002	Alice	James	123457	07/31/1972	F	980 Berry Ln, Elgin,IL
3 E1003	Steve	Wells	123458	08/10/1980	M	291 Springs, Gary,IL
4 E1004	Santosh	Kumar	123459	07/20/1985	M	511 Aurora Av, Aurora,IL
5 E1005	Ahmed	Hussain	123410	01/04/1981	M	216 Oak Tree, Geneva,IL
6 E1006	Nancy	Alien	123411	02/06/1978	F	111 Green Pl, Elgin,IL
7 E1007	Mary	Thomas	123412	05/05/1975	F	100 Rose Pl, Gary,IL
8 E1008	Bharath	Gupta	123413	05/06/1985	M	145 Berry Ln, Naperville,IL
9 E1009	Andrea	Jones	123414	07/09/1990	F	120 Fall Creek, Gary,IL

- 7) Click "Next". Review the Load setting and click "Begin Load" in the top-right corner.

Review settings

Summary

Code page:	1208 <i>(Default)</i>
Separator:	,
Header in first row:	No
Time format:	HH:MM:SS <i>(Default)</i>
Date format:	MM/DD/YYYY
Timestamp format:	YYYY-MM-DD HH:MM:SS <i>(Default)</i>
String delimiter:	" <i>(Default)</i>

Option

Maximum number of warnings

1000

[Back](#)
Begin Load

- 8) After Loading is complete you will notice that we were successful in loading all 10 rows of the Employees table. If there are any Errors or Warnings you can view them on this screen.

Load details

✓
COMPLETE

My computer Target

Employees.csv QCM54853.EMPLOYEES

View Table
Load More Data

Status [Settings](#)

10 10 0

Rows read Rows loaded Rows rejected

Start time: 05/02/2018 1:51:27 PM
End time: 05/02/2018 1:51:28 PM

The data load job succeeded.
You can now work with your data.

No errors

- 9) You can see the data that was loaded by clicking on the View Table. Alternatively, you can go into the Explore page and page select the correct schema, then the EMPLOYEES table, and click "View Data".

QCM54853.EMPLOYEES

Delete Table Export to CSV

	EMP_ID CHARACTER(9)	F_NAME VARCHAR(15)	L_NAME VARCHAR(15)	SSN CHARACTER(9)	B_DATE DATE	SEX CHARACTER(1)	ADDRESS VARCHAR(30)	JOB_ID CHARACTER(9)
1	E1001	John	Thomas	123456	1976-01-09	M	5631 Rice, OakPark, 100	
2	E1002	Alice	James	123457	1972-07-31	F	980 Berry Ln, Elgin,IL 200	
3	E1003	Steve	Wells	123458	1980-08-10	M	291 Springs, Gary,IL 300	
4	E1004	Santosh	Kumar	123459	1985-07-20	M	511 Aurora Av, Aurora, IL 400	
5	E1005	Ahmed	Hussain	123410	1981-01-04	M	216 Oak Tree, Geneva, IL 500	
6	E1006	Nancy	Allen	123411	1978-02-06	F	111 Green Pl, Elgin,IL 600	
7	E1007	Mary	Thomas	123412	1975-05-05	F	100 Rose Pl, Gary,IL 650	
8	E1008	Bharath	Gupta	123413	1985-05-06	M	145 Berry Ln, Naper, IL 660	
9	E1009	Andrea	Jones	123414	1990-07-09	F	120 Fall Creek, Gary, IL 234	
10	E1010	Ann	Jacob	123415	1982-03-30	F	111 Britany Springs, Elgin, IL 220	

10. Now its your turn to load the remaining 4 tables of the HR database – Locations, JobHistory, Jobs, and Departments. Please follow the steps above to load the data from the remaining source files.

Question 1: Were there any warnings loading data into the JOBS table? What can be done to resolve this?

Hint: View the data loaded into this table and pay close attention to the JOB_TITLE column.

Question 2: Did all rows from the source file load successfully in the DEPARTMENT table? If not, are you able to figure out why not?

Hint: Look at the warning. Also, note the Primary Key for this table.

Part III: COMPOSING AND RUNNING QUERIES

You created the tables for the HR database schema and also learned how to load data into these tables. Now try and work on a few advanced DML queries that were introduced in this module.

Follow these steps to create and run the queries indicated below

- 1) Navigate to the Run SQL tool in Db2 on Cloud
- 2) Compose query and run it.
- 3) Check the Logs created under the Results section. Looking at the contents of the Log explains whether the SQL statement ran successfully. Also look at the query results to ensure the output is what you expected.

Query 1: Retrieve all employees whose address is in Elgin,IL

[Hint: Use the LIKE operator to find similar strings]

Query 2: Retrieve all employees who were born during the 1970's.

[Hint: Use the LIKE operator to find similar strings]

Query 3: Retrieve all employees in department 5 whose salary is between 60000 and 70000 .

[Hint: Use the keyword BETWEEN for this query]

Query 4A: Retrieve a list of employees ordered by department ID.

[Hint: Use the ORDER BY clause for this query]

Query 4B: Retrieve a list of employees ordered in descending order by department ID and within each department ordered alphabetically in descending order by last name.

Query 5A: For each department ID retrieve the number of employees in the department.

[Hint: Use COUNT(*) to retrieve the total count of a column, and then GROUP BY]

Query 5B: For each department retrieve the number of employees in the department, and the average employees salary in the department.

[Hint: Use COUNT(*) to retrieve the total count of a column, and AVG() function to compute average salaries, and then group]

Query 5C: Label the computed columns in the result set of Query 5B as “NUM_EMPLOYEES” and “AVG_SALARY”.

[Hint: Use AS “LABEL_NAME” after the column name]

Query 5D: In Query 5C order the result set by Average Salary.

[Hint: Use ORDER BY after the GROUP BY]

Query 5E: In Query 5D limit the result to departments with fewer than 4 employees.

[Hint: Use HAVING after the GROUP BY, and use the count() function in the HAVING clause instead of the column label.

Note: WHERE clause is used for filtering the entire result set whereas the HAVING clause is used for filtering the result of the grouping]

BONUS Query 6: Similar to 4B but instead of department ID use department name. Retrieve a list of employees ordered by department name, and within each department ordered alphabetically in descending order by last name.

[Hint: Department name is in the DEPARTMENTS table. So your query will need to retrieve data from more than one table. Don't worry if you are not able to figure this one out ... we'll cover working with multiple tables in the next lesson.]

In this lab you learned how to work with string patterns, sorting result sets and grouping result sets.

Thank you for completing this lab!

See solutions on the following page.

Lab Solutions

Please follow these steps to get the answers to the queries:

- 1) Navigate to the Run SQL page on Db2 on Cloud.
- 2) Download the script file("Module4_Qualities.txt") or text files("Module4_Qualities.sql"). Open the file with extension .sql in the editor
- 3) Run the queries. Looking at the contents of the Log explains that the SQL statement that we ran was successful. Here are the results for the queries:

Query 1: Output

```

1 -- Query 1-----
2 ;
3 select F_NAME , L_NAME
4 from EMPLOYEES
5 where ADDRESS LIKE '%Elgin,IL%' ;
6 --Query 2--
7 ;

```

Saved scripts	Result
Filter by status: Result set Log	All
Delete All	F_NAME L_NAME
Al... 	Alice James
select...	Nancy Allen
select ...	Ann Jacob
select ...	Total rows: 3

Query 2: Output

RUN SQL

Run ▾ Script ▾ Edit ▾ Favorites ▾ New tab

```

6 --Query 2--|
7 ;
8 select F_NAME , L_NAME
9 from EMPLOYEES
10 where B_DATE LIKE '197%';
11 ---Query3---
12 :

```

Saved scripts	Result
Filter by status: Result set Log	All
Delete All	F_NAME L_NAME
Al... 	John Thomas
select...	Alice James
select ...	Nancy Allen
select ...	Mary Thomas
select ...	Total rows: 4

Query 3: Output

```

11 ---Query3---|
12 ;
13 select *
14 from EMPLOYEES
15 where (SALARY BETWEEN 60000 and 70000) and DEP_ID = 5 ;
16 --Query4--|
17 ;

```

Saved scripts Result

Filter by status: Result set Log

All

Delete All

- ✓ All(5)...
- ✓ select F....
- ✓ select F....
- ✓ select * f...

EMP_ID F_NAME L_NAME SSN B_DATE SEX ADDRESS JOB_ID SALARY MANAGER_ID DEP_ID

E1004	Santosh	Kumar	1234...	1985-07-20	M	511 Aurora ...	400	60000.00	30004	5
E1010	Ann	Jacob	12341...	1982-03-30	F	111 Britany ...	220	70000.00	30004	5

Total rows: 2

Query 4A: Output

```

select F_NAME, L_NAME, DEP_ID
from EMPLOYEES
order by DEP_ID;

```

Saved scripts Result

Filter by status: Result set Log

All

Delete All

- ✓ (1)...
- ✓ select F....

F_NAME L_NAME DEP_ID

John	Thomas	2
Ahmed	Hussain	2
Nancy	Allen	2
Alice	James	5
Steve	Wells	5
Santosh	Kumar	5
Ann	Jacob	5
Mary	Thomas	7
Bharath	Gupta	7
Andrea	Jones	7

Total rows: 10

Query 4B: Output

```
1 select F_NAME, L_NAME, DEP_ID
2 from EMPLOYEES
3 order by DEP_ID desc, L_NAME desc;
```

Saved scripts	Result
Filter by status:	
	Result set
All	Log
Delete All	
✓ All(1)...	
(select F...	
✓ All(1)...	
(select F...	
✓ All(1)...	
(select F...	
✓ All(1)...	
(select F...	
✓ All(1)...	
(select F...	
Total rows: 10	

Query 5A: Output

```
select DEP_ID, COUNT(*)
from EMPLOYEES
group by DEP_ID;
```

Saved scripts	Result
Filter by status:	
All	Result set
Delete All	
J...	
(select...	
J...	
(select ...	
Total rows: 3	

Query 5B: Output

```

1 select DEP_ID, COUNT(*), AVG(SALARY)
2 from EMPLOYEES
3 group by DEP_ID;

```

DEP_ID	Count	Avg_Salary
2	3	86666.6666666666666666666666666666
5	4	65000.0000000000000000000000000000
7	3	66666.6666666666666666666666666666

Total rows: 3

Query 5C: Output

```

select DEP_ID, COUNT(*) AS "NUM_EMPLOYEES", AVG(SALARY) AS "AVG_SALARY"
from EMPLOYEES
group by DEP_ID;

```

DEP_ID	NUM_EMPLOYEES	AVG_SALARY
2	3	86666.6666666666666666666666666666
5	4	65000.0000000000000000000000000000
7	3	66666.6666666666666666666666666666

Total rows: 3

Query 5D: Output

```

select DEP_ID, COUNT(*) AS "NUM_EMPLOYEES", AVG(SALARY) AS "AVG_SALARY"
from EMPLOYEES
group by DEP_ID
order by AVG_SALARY;

```

DEP_ID	NUM_EMPLOYEES	AVG_SALARY
5	4	65000.0000000000000000000000000000
7	3	66666.6666666666666666666666666666
2	3	86666.6666666666666666666666666666

Total rows: 3

Query 5E: Output

```
select DEP_ID, COUNT(*) AS "NUM_EMPLOYEES", AVG(SALARY) AS "AVG_SALARY"
from EMPLOYEES
group by DEP_ID
having count(*) < 4
order by AVG_SALARY;
```

d scripts **Result**

by status: **Result set** Log

Delete All

!& F...

DEP_ID	NUM_EMPLOYEES
7	3 66666.6666666
2	3 86666.6666666

Total rows: 2

BONUS Query 6: Output

Note that in the Query below “D” and “E” are aliases for the table names. Once you define an alias like “D’ in your query, you can simply write “D.COLUMN_NAME” rather than the full form ‘DEPARTMENTS.COLUMN_NAME’.

```
16 --Query4--
17 ;
18 select D.DEP_NAME , E.F_NAME, E.L_NAME
19 from EMPLOYEES as E, DEPARTMENTS as D
20 where E.DEP_ID = D.DEP_ID
21 order by D.DEP_NAME, E.L_NAME desc ;
22 --Query5--
```

Saved scripts **Result**

Filter by status: **Result set** Log

Delete All

✓ All(5)...

✓ select F_...
✓ select F_...
✓ select * fr...
✓ select D...
✓ select DE...

DEP_NAME	F_NAME	L_NAME
Architect Group	John	Thomas
Architect Group	Ahmed	Hussain
Architect Group	Nancy	Allen
Design Team	Mary	Thomas
Design Team	Andrea	Jones
Design Team	Bharath	Gupta
Software Group	Steve	Wells
Software Group	Santosh	Kumar
Software Group	Alice	James
Software Group	Ann	Jacob

Total rows: 10